




Attention is the mixing step (tokens talk to each other).
The FFN is the thinking step (each token computes alone).
Cost grows with sequence length — that is why context is finite.